

CS 33

Data Representation (Part 2)

Why Should I Use Unsigned?

- ***Don't*** use just because number nonnegative

- easy to make mistakes

```
unsigned i;  
for (i = cnt-2; i >= 0; i--)  
    a[i] += a[i+1];
```

- can be very subtle

```
#define DELTA sizeof(int)  
int i;  
for (i = CNT; i-DELTA >= 0; i-= DELTA)  
    . . .
```

- ***Do*** use when using bits to represent sets
 - logical right shift, no sign extension

Word Size

- **(Mostly) obsolete term**
 - old computers had items of one size: the word size
- **Now used to express the number of bits necessary to hold an address**
 - 16 bits (really old computers)
 - 32 bits (old computers)
 - 64 bits (most current computers)

Byte Ordering

- **Four-byte integer**
 - 0x76543210
- **Stored at location 0x100**
 - which byte is at 0x100?
 - which byte is at 0x103?

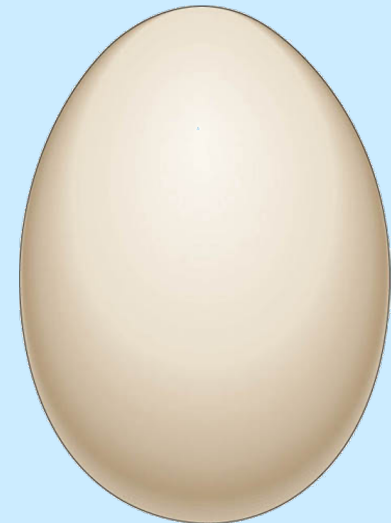


10	32	54	76
0x100	0x101	0x102	0x103

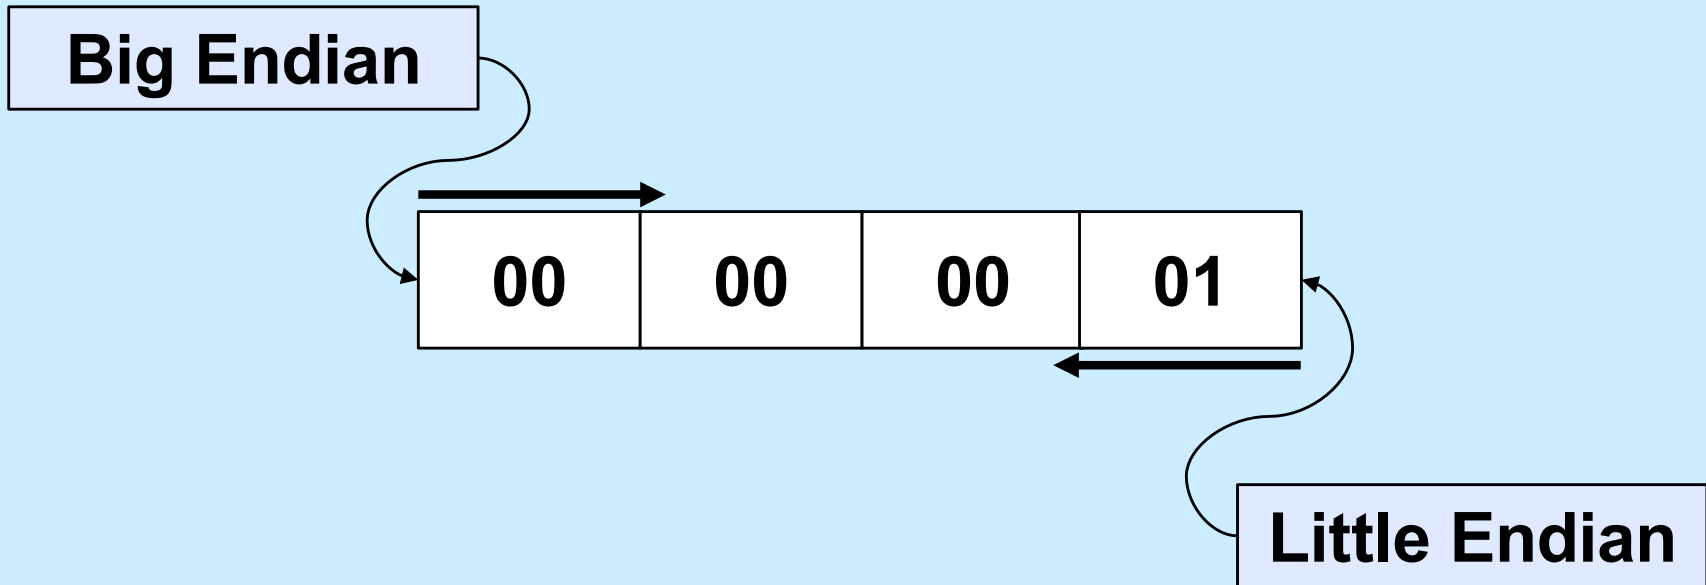
Little-endian

76	54	32	10
0x100	0x101	0x102	0x103

Big-endian



Byte Ordering (2)



Quiz 1

```
int main() {  
    long x=1;  
    func((int *) &x);  
    return 0;  
}  
  
void func(int *arg) {  
    printf("%d\n", *arg);  
}
```

**What value is printed
on a big-endian 64-bit
computer?**

- a) 1
- b) 0
- c) 2^{32}
- d) $2^{32}-1$

Which Byte Ordering Do We Use?

```
int main() {  
    unsigned int x = 0x03020100;  
    unsigned char *xarray = (unsigned char *) &x;  
    for (int i=0; i<4; i++) {  
        printf("%02x", xarray[i]);  
    }  
    printf("\n");  
    return 0;  
}
```

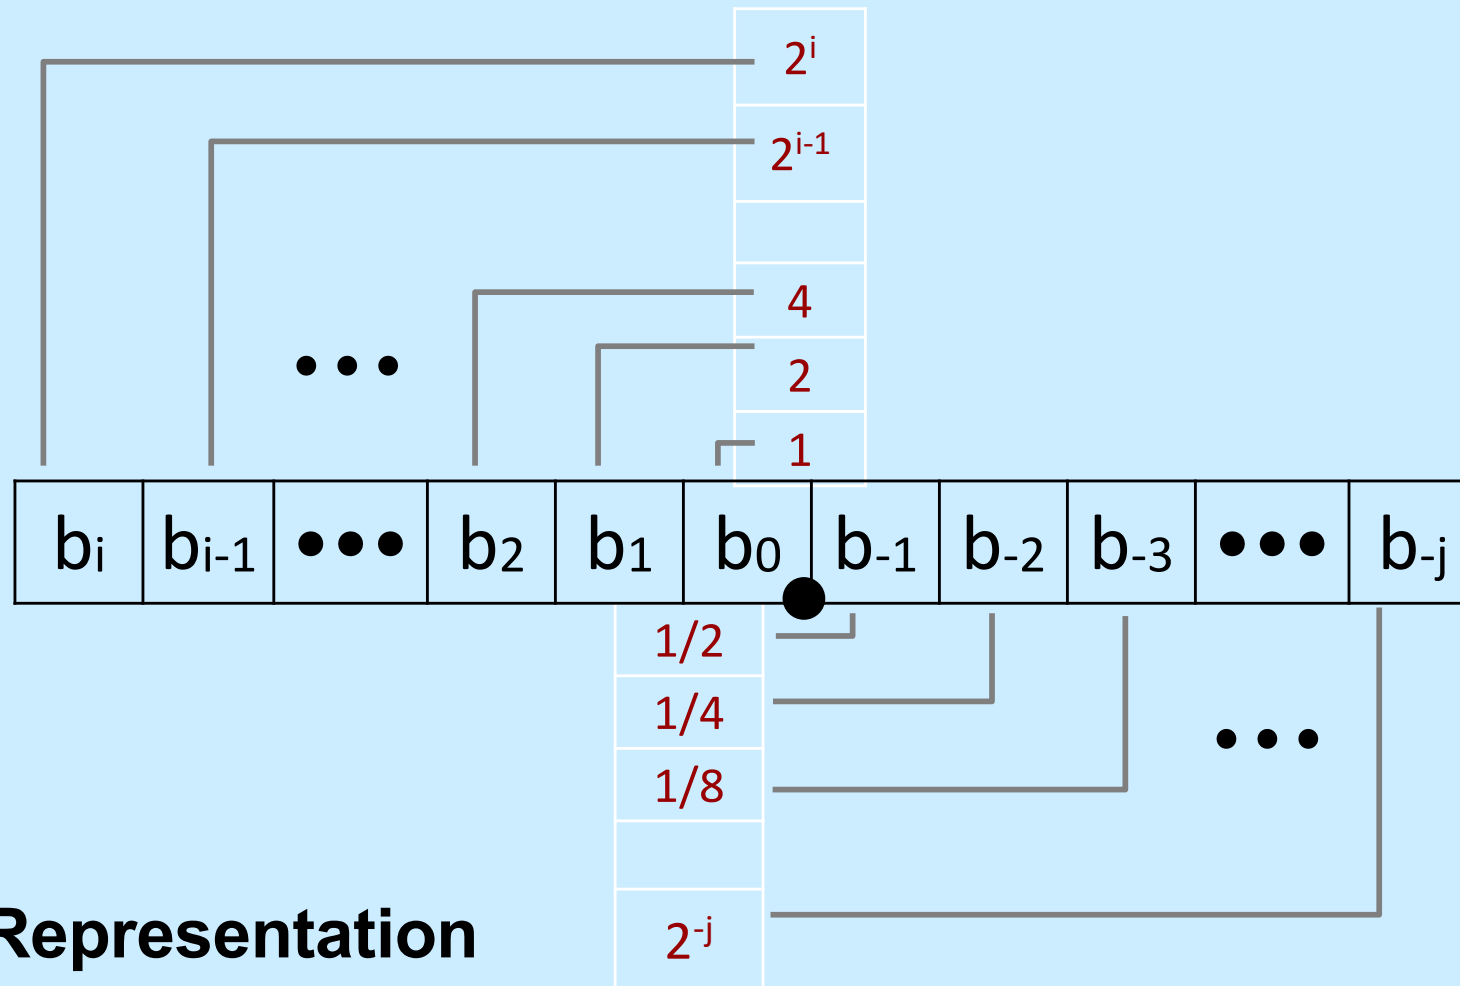
Possible results:

00010203
03020100

Fractional binary numbers

- What is 1011.101_2 ?

Fractional Binary Numbers



- **Representation**

- bits to right of “binary point” represent fractional powers of 2
- represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$

Representable Numbers

- **Limitation #1**

- can exactly represent only numbers of the form $n/2^k$

- » other rational numbers have repeating bit representations

- value representation

- » 1/3 0.0101010101[01]...₂

- » 1/5 0.001100110011[0011]...₂

- » 1/10 0.0001100110011[0011]...₂

- **Limitation #2**

- just one setting of decimal point within the w bits

- » limited range of numbers (very small values? very large?)

IEEE Floating Point

- **IEEE Standard 754**
 - established in 1985 as uniform standard for floating point arithmetic
 - » before that, many idiosyncratic formats
 - supported on all major CPUs
- **Driven by numerical concerns**
 - nice standards for rounding, overflow, underflow
 - hard to make fast in hardware
 - » numerical analysts predominated over hardware designers in defining standard

Floating-Point Representation

- Numerical Form:

$$(-1)^s M 2^E$$

- sign bit **s** determines whether number is negative or positive
- significand **M** normally a fractional value in range $[1.0, 2.0)$
- exponent **E** weights value by power of two

- Encoding

- MSB **s** is sign bit **s**
- exp field encodes **E** (but is not equal to E)
- frac field encodes **M** (but is not equal to M)



Precision options

- **Single precision: 32 bits**



- **Double precision: 64 bits**



- **Extended precision: 80 bits (Intel only)**



“Normalized” Values

- When: $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$
- Exponent coded as biased value: $E = \text{Exp} - \text{Bias}$
 - exp : unsigned value exp
 - $\text{bias} = 2^{k-1} - 1$, where k is number of exponent bits
 - » single precision: 127 (Exp: 1...254, E: -126...127)
 - » double precision: 1023 (Exp: 1...2046, E: -1022...1023)
- Significand coded with implied leading 1: $M = 1.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
 - minimum when $\text{frac} = 000\dots 0$ ($M = 1.0$)
 - maximum when $\text{frac} = 111\dots 1$ ($M = 2.0 - \epsilon$)
 - get extra leading bit for “free”

Normalized Encoding Example

- Value: `float F = 15213.0;`

$$\begin{aligned} - 15213_{10} &= 11101101101101_2 \\ &= 1.1101101101101_2 \times 2^{13} \end{aligned}$$

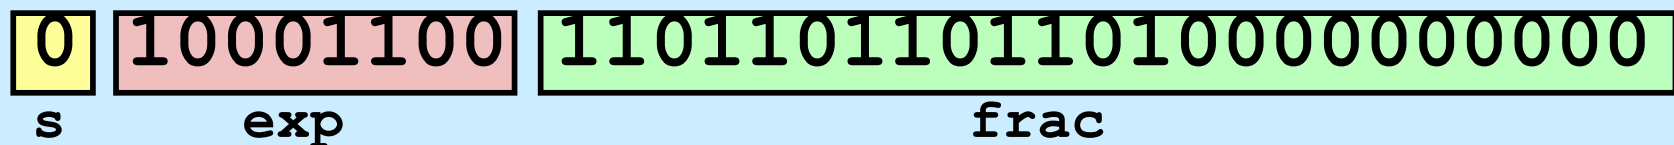
- Significand

$$\begin{aligned} M &= 1.\underline{1101101101101}_2 \\ \text{frac} &= \underline{1101101101101}0000000000_2 \end{aligned}$$

- Exponent

$$\begin{aligned} E &= 13 \\ \text{bias} &= 127 \\ \text{exp} &= 140 = 10001100_2 \end{aligned}$$

- Result:



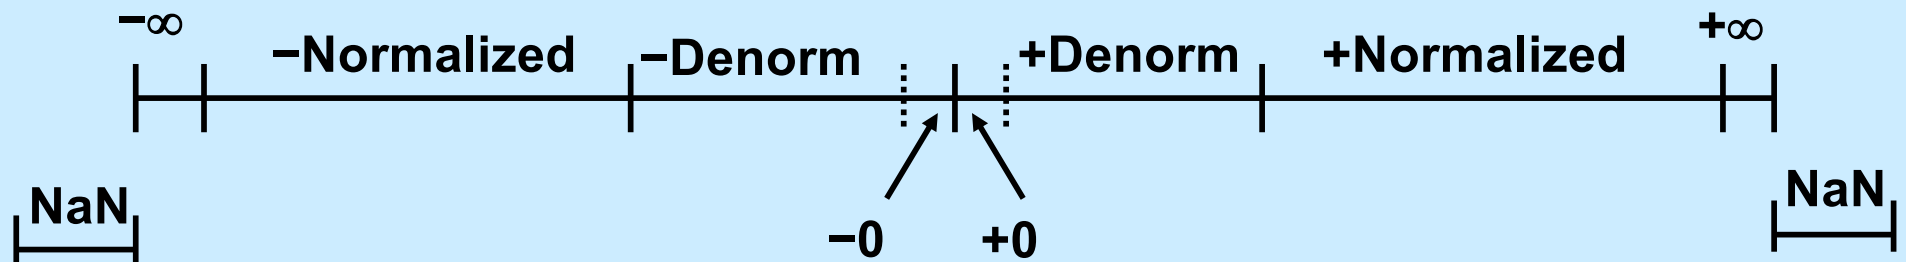
Denormalized Values

- **Condition:** $\text{exp} = 000\dots 0$
- **Exponent value:** $E = -\text{Bias} + 1$ (instead of $E = 0 - \text{Bias}$)
- **Significand coded with implied leading 0:**
 $M = 0.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
- **Cases**
 - $\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$
 - » represents zero value
 - » note distinct values: $+0$ and -0 (why?)
 - $\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$
 - » numbers closest to 0.0
 - » equispaced

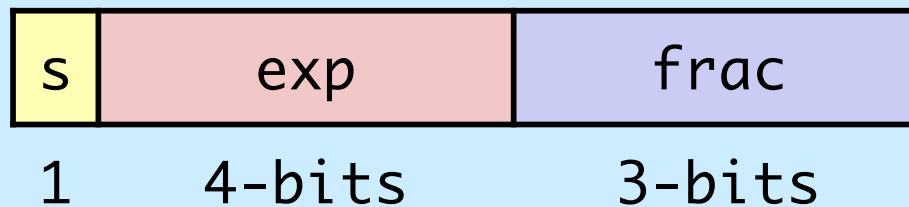
Special Values

- **Condition: $\text{exp} = 111\dots 1$**
 - **Case: $\text{exp} = 111\dots 1$, $\text{frac} = 000\dots 0$**
 - represents value ∞ (infinity)
 - operation that overflows
 - both positive and negative
 - e.g., $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$
 - **Case: $\text{exp} = 111\dots 1$, $\text{frac} \neq 000\dots 0$**
 - not-a-number (NaN)
 - represents case when no numeric value can be determined
 - e.g., $\text{sqrt}(-1)$, $\infty - \infty$, $\infty \times 0$
-

Visualization: Floating-Point Encodings



Tiny Floating-Point Example



- **8-bit Floating Point Representation**
 - the sign bit is in the most significant bit
 - the next four bits are the exponent, with a bias of 7
 - the last three bits are the *frac*
- **Same general form as IEEE Format**
 - normalized, denormalized
 - representation of 0, NaN, infinity

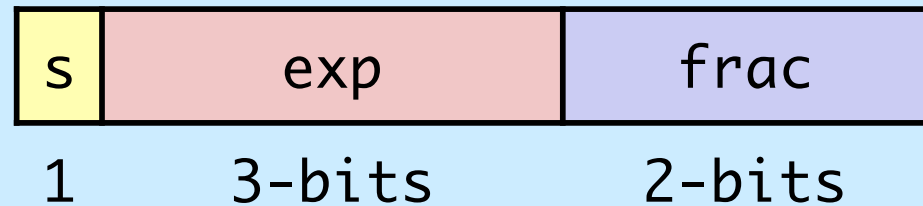
Dynamic Range (Positive Only)

	s	exp	frac	E	Value	
Denormalized numbers	0	0000	000	-6	0	
	0	0000	001	-6	$1/8 * 1/64 = 1/512$	closest to zero
	0	0000	010	-6	$2/8 * 1/64 = 2/512$	
	...					
	0	0000	110	-6	$6/8 * 1/64 = 6/512$	
	0	0000	111	-6	$7/8 * 1/64 = 7/512$	largest denorm
Normalized numbers	0	0001	000	-6	$8/8 * 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 * 1/64 = 9/512$	
	...					
	0	0110	110	-1	$14/8 * 1/2 = 14/16$	
	0	0110	111	-1	$15/8 * 1/2 = 15/16$	closest to 1 below
	0	0111	000	0	$8/8 * 1 = 1$	
	0	0111	001	0	$9/8 * 1 = 9/8$	closest to 1 above
	0	0111	010	0	$10/8 * 1 = 10/8$	
	...					
	0	1110	110	7	$14/8 * 128 = 224$	
	0	1110	111	7	$15/8 * 128 = 240$	largest norm
	0	1111	000	n/a	inf	

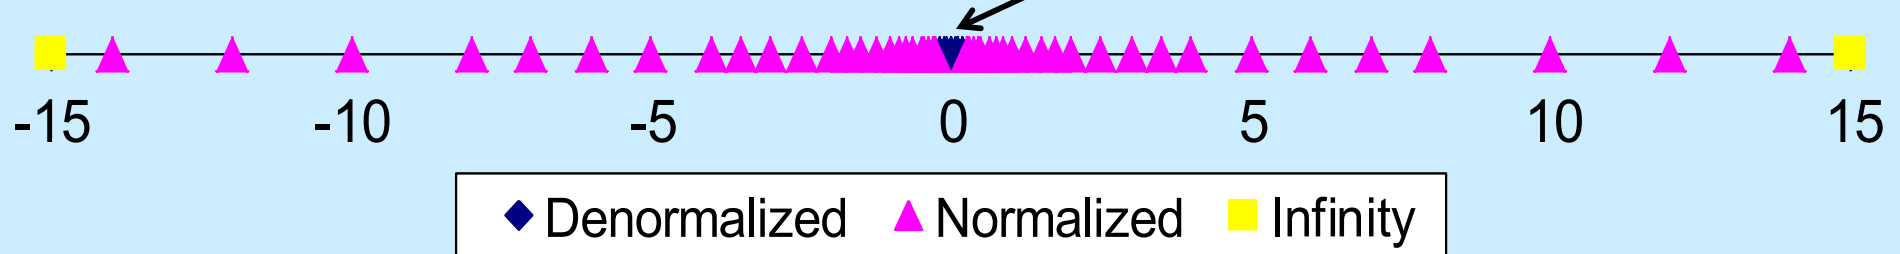
Distribution of Values

- 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- bias is $2^{3-1}-1 = 3$



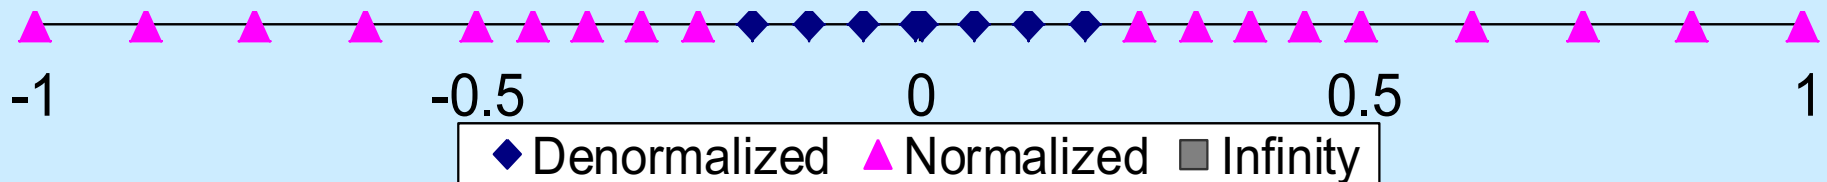
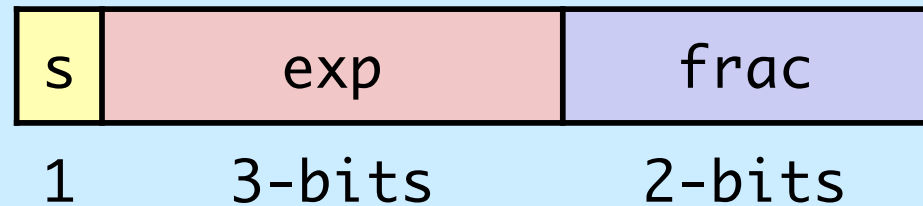
- Notice how the distribution gets denser toward zero.



Distribution of Values (close-up view)

- 6-bit IEEE-like format

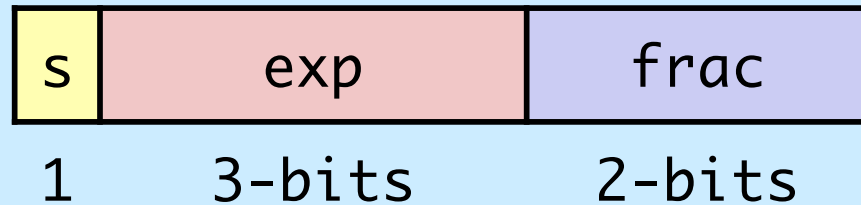
- e = 3 exponent bits
- f = 2 fraction bits
- bias is 3



Quiz 1

- **6-bit IEEE-like format**

- e = 3 exponent bits
- f = 2 fraction bits
- bias is 3

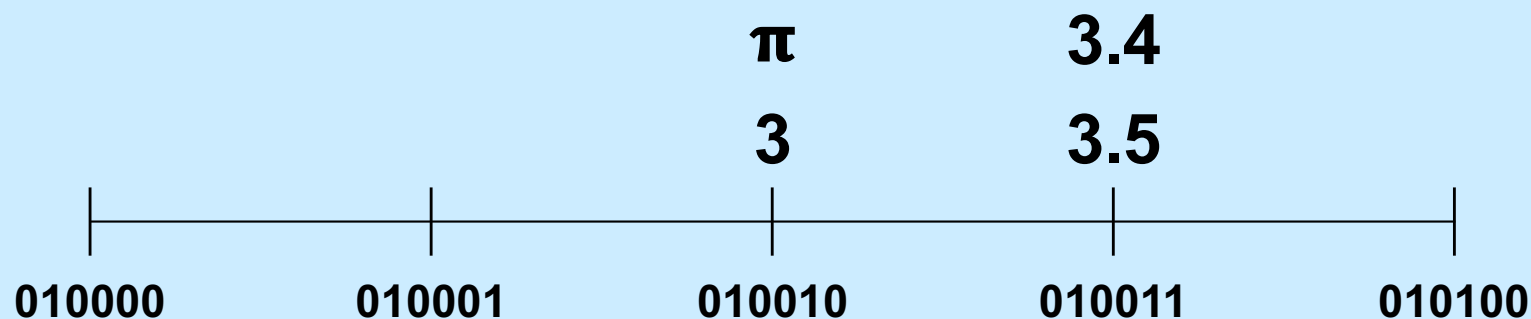


What number is represented by 0 010 10?

- a) 3
- b) 1.5
- c) .75
- d) none of the above

Mapping Real Numbers to Float

- The real number 3 is represented as
0 100 10
- The real number 3.5 is represented as
0 100 11
- How is the real number 3.4 represented?
0 100 11
- How is the real number π represented?
0 100 10

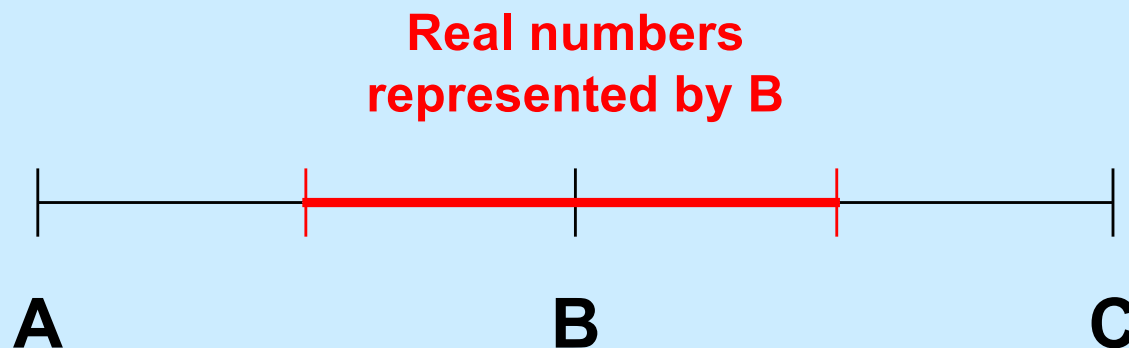


Mapping Real Numbers to Float

- If R is a real number, it's mapped to the floating-point number whose value is closest to R
- What if it's midway between two values?
 - rounding rules coming up soon!

Floats are Sets of Values

- If A, B, and C are successive floating-point numbers
 - e.g., 010001, 010010, and 010011
- B represents all real numbers from midway between A and B through midway between B and C



Significance

- **Normalized numbers**
 - for a particular exponent value E and an S -bit significand, the range from 2^E up to 2^{E+1} is divided into 2^S equi-spaced floating-point values
 - » thus each floating-point value represents $1/2^S$ of the range of values with that exponent
 - » all bits of the significand are important
 - » we say that there are S significant bits – for reasonably large S , each floating-point value covers a rather small part of the range
 - high accuracy
 - for $S=23$ (32-bit float), accurate to one in 2^{23} (.0000119% accuracy)

Significance

- **Unnormalized numbers**
 - high-order zero bits of the significand aren't important
 - in 8-bit floating point, 0 0000 001 represents 2^{-9}
 - » it is the only value with that exponent: 1 significant bit (either 2^{-9} or 0)
 - 0 0000 010 represents 2^{-8}
0 0000 011 represents $1.5 \cdot 2^{-8}$
 - » only two values with exponent -8: 2 significant bits (encoding those two values, as well as 2^{-9} and 0)
 - fewer significant bits mean less accuracy
 - 0 0000 001 represents a range of values from $.5 \cdot 2^{-9}$ to $1.5 \cdot 2^{-9}$
 - 50% accuracy

Floating-Point Operations: Basic Idea

- $x +_f y = \text{Round}(x + y)$
- $x \times_f y = \text{Round}(x \times y)$
- Basic idea
 - first **compute exact result**
 - make it fit into desired precision
 - » possibly overflow if exponent too large
 - » possibly **round to fit into** frac

Rounding

- Rounding modes (illustrated with \$ rounding)

	\$1.40	\$1.60	\$1.50	\$2.50	−\$1.50
towards zero	\$1	\$1	\$1	\$2	−\$1
round down ($-\infty$)	\$1	\$1	\$1	\$2	−\$2
round up ($+\infty$)	\$2	\$2	\$2	\$3	−\$1
nearest integer	\$1	\$2	?	?	?
nearest even (default)	\$1	\$2	\$2	\$2	−\$2

Floating-Point Multiplication

- $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$
- **Exact result:** $(-1)^s M 2^E$
 - sign s : $s1 \wedge s2$
 - significand M : $M1 \times M2$
 - exponent E : $E1 + E2$
- **Fixing**
 - if $M \geq 2$, shift M right, increment E
 - if E out of range, overflow (or underflow)
 - round M to fit `frac` precision
- **Implementation**
 - biggest chore is multiplying significands

Floating-Point Addition

- $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

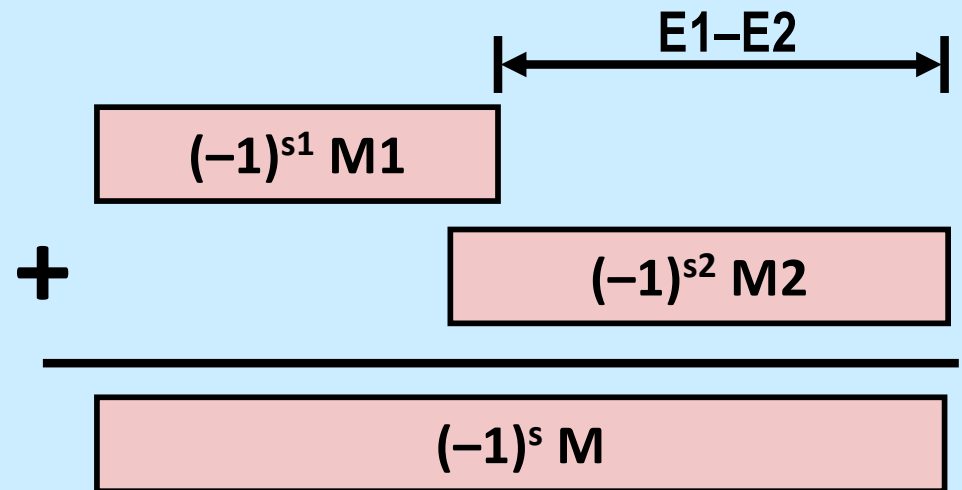
–assume $E1 > E2$

- **Exact result:** $(-1)^s M 2^E$

–sign s , significand M :

» result of signed align & add

–exponent E : $E1$



- **Fixing**

–if $M \geq 2$, shift M right, increment E

–if $M < 1$, shift M left k positions, decrement E by k

–overflow if E out of range

–round M to fit **frac** precision