

# CS 33

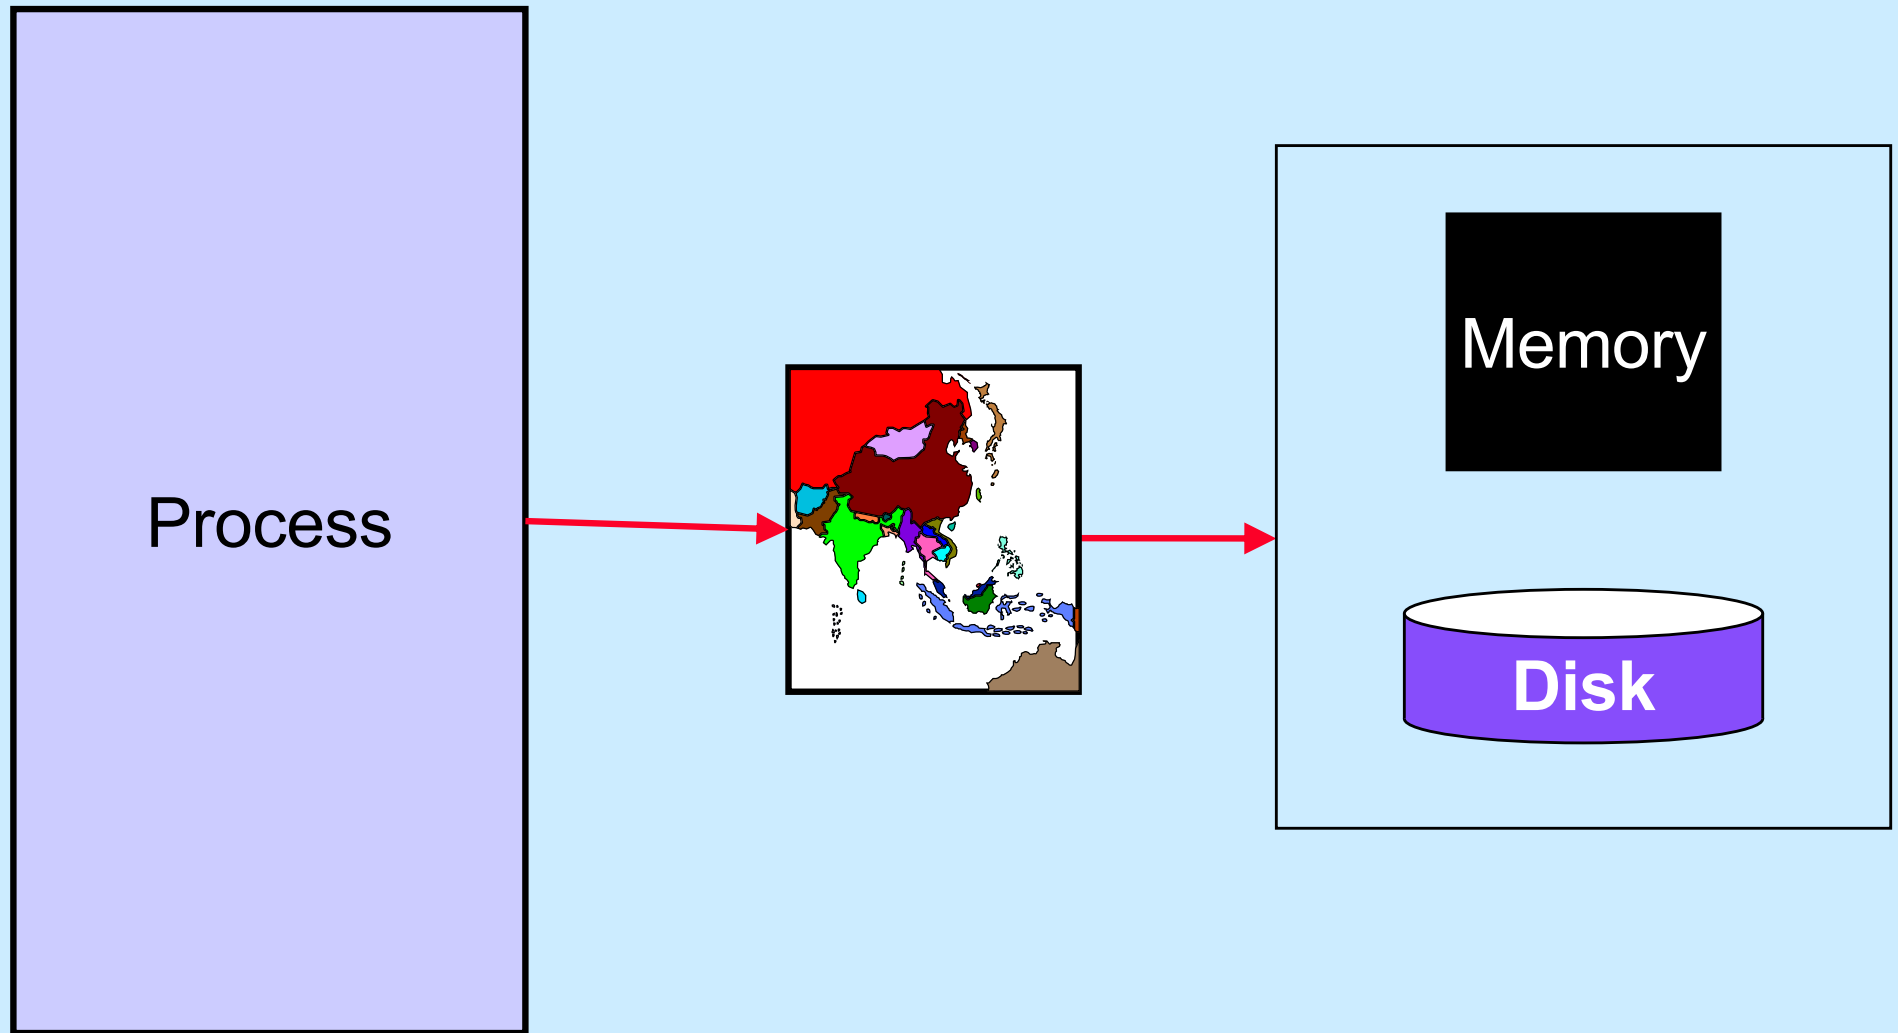
## Virtual Memory 2

# OS Role in Virtual Memory


- **Memory is like a cache**
  - quick access if what's wanted is mapped via page table
  - slow if not — OS assistance required
- **OS**
  - make sure what's needed is mapped in
  - make sure what's no longer needed is not mapped in

# Why is virtual memory used?

# More VM than RM




# Isolation






Virtual Memory

# Sharing




**Virtual Memory**

# File I/O




# Multi-Buffered I/O








# Traditional I/O



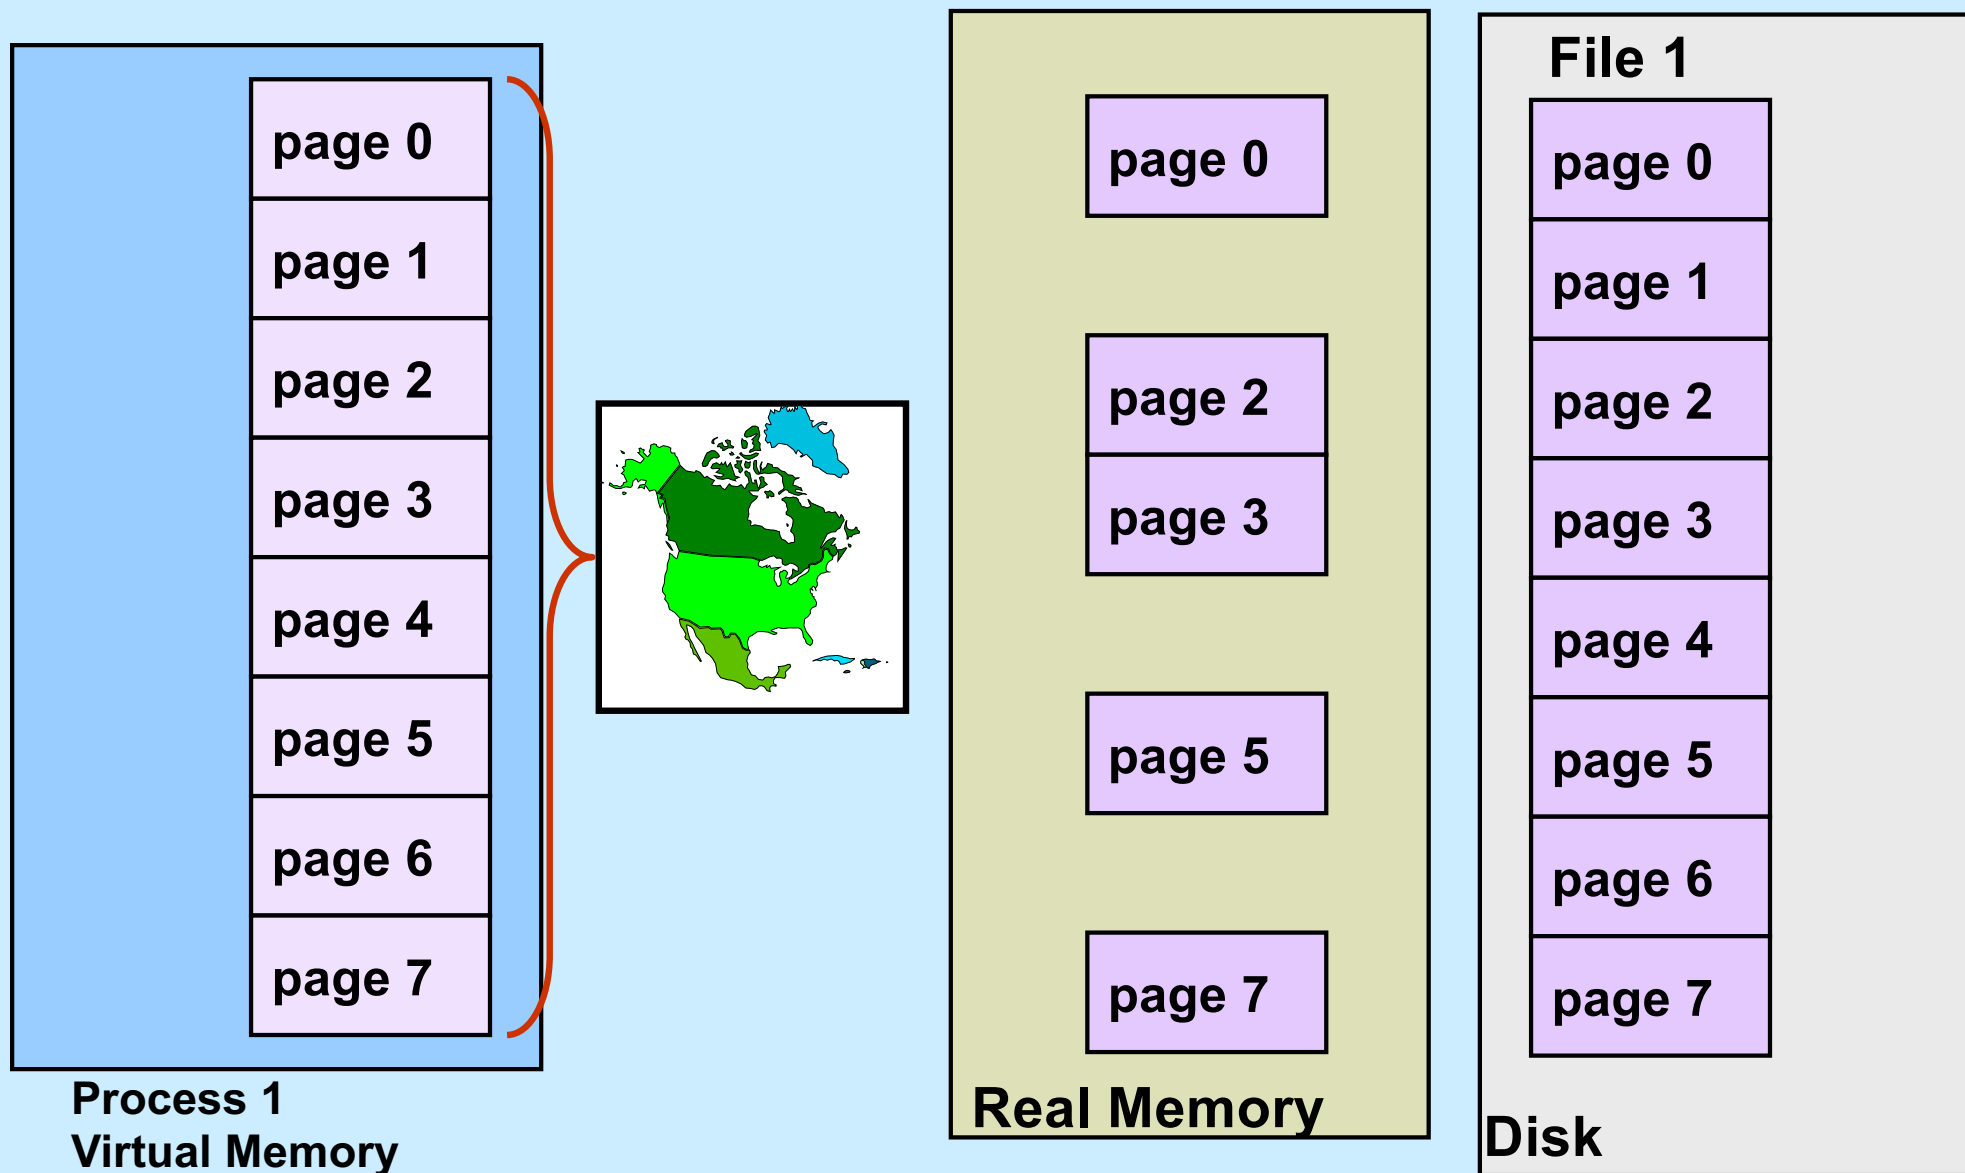
User Process 1



User Process 2



# Mapped File I/O

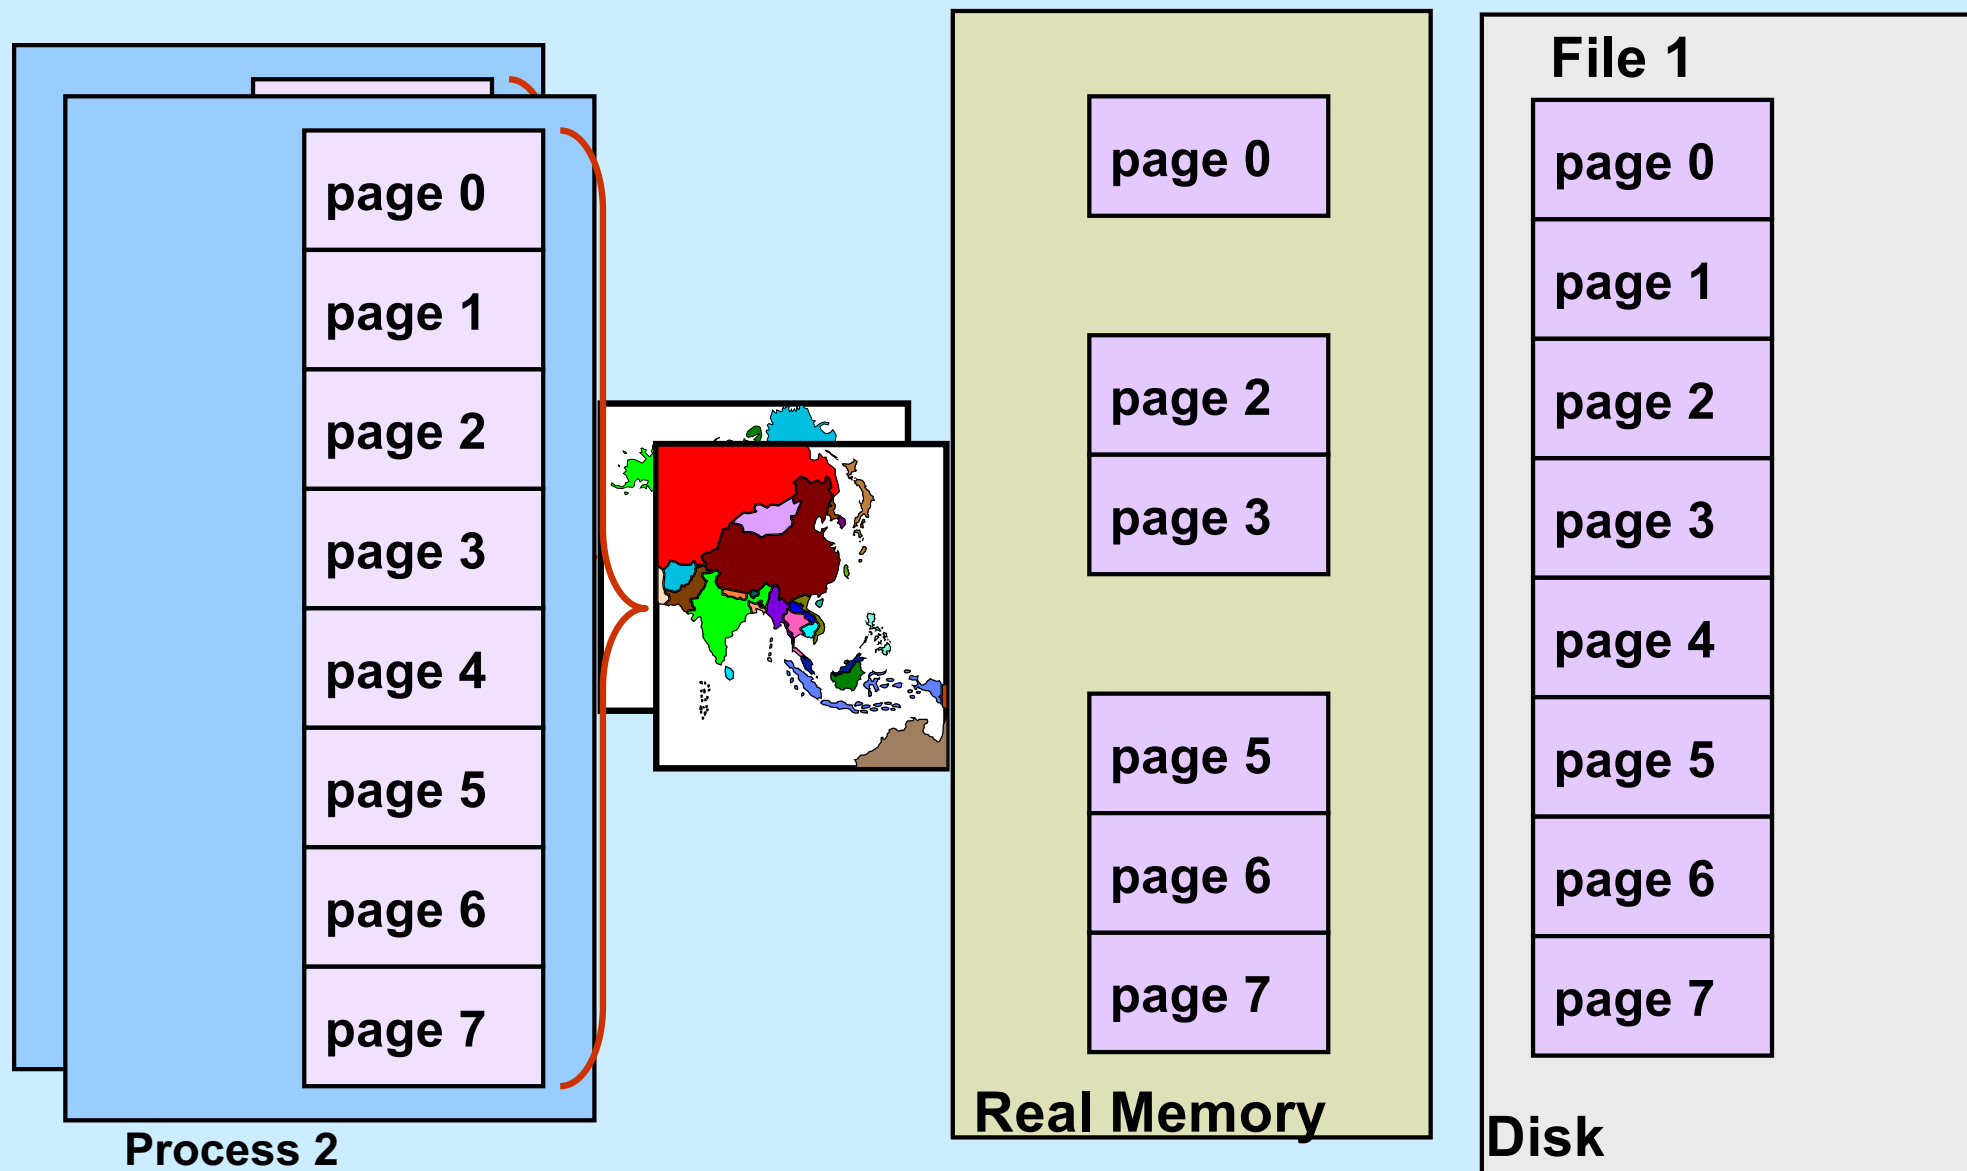


Process 1  
Virtual Memory

Real Memory

Disk

# Multi-Process Mapped File I/O



# Mapped Files

- **Traditional File I/O**

```
char buf[BigEnough];  
fd = open(file, O_RDWR);  
for (i=0; i<n_recs; i++) {  
    read(fd, buf, sizeof(buf));  
    use(buf);  
}
```


- **Mapped File I/O**

```
record_t *MappedFile;  
fd = open(file, O_RDWR);  
MappedFile = mmap(... , fd, ...);  
for (i=0; i<n_recs; i++)  
    use(MappedFile[i]);
```


# Mmap System Call

```
void *mmap(  
    void *addr,  
        // where to map file (0 if don't care)  
    size_t len,  
        // how much to map  
    int prot,  
        // memory protection (read, write, exec.)  
    int flags,  
        // shared vs. private, plus more  
    int fd,  
        // which file  
    off_t off  
        // starting from where  
);
```

# The *mmap* System Call




# Share-Mapped Files



Data = 17;

# Private-Mapped Files





# Example

```
int main( ) {
    int fd;
    dataObject_t *dataObjectp;

    fd = open("file", O_RDWR);
    if ((int)(dataObjectp = (dataObject_t *)mmap(0,
        sizeof(dataObject_t),
        PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0)) == -1) {
        perror("mmap");
        exit(1);
    }

    // dataObjectp points to region of (virtual) memory
    // containing the contents of the file

    ...

}
```

# fork and mmap

```
int main() {
    int x=1;

    if (fork() == 0) {
        // in child
        x = 2;
        exit(0);
    }
    // in parent
    while (x==1) {
        // will loop forever
    }
    return 0;
}
```

```
int main() {
    int fd = open( ... );
    int *xp = (int *)mmap(...,
        MAP_SHARED, fd, ...);
    xp[0] = 1;
    if (fork() == 0) {
        // in child
        xp[0] = 2;
        exit(0);
    }
    // in parent
    while (xp[0]==1) {
        // will terminate
    }
    return 0;
}
```

# Putting Together a Program

# gcc Steps

## 1) Compile

- to start here, supply .c file
- to stop here: `gcc -S` (produces .s file)
- if not stopping here, gcc compiles directly into a .o file, bypassing the assembler

## 2) Assemble

- to start here, supply .s file
- to stop here: `gcc -c` (produces .o file)

## 3) Link

- to start here, supply .o file

# The Linker

- **An executable program is one that is ready to be loaded into memory**
- **The linker (known as ld: /usr/bin/ld) creates such executables from:**
  - **object files produced by the compiler/assembler**
  - **collections of object files (known as libraries or archives)**
  - **and more we'll get to soon ...**

# Linker's Job

- **Piece together components of program**
  - **arrange within address space**
    - » **code (and read-only data) goes into text region**
    - » **initialized data goes into data region**
    - » **uninitialized data goes into bss region**
- **Modify address references, as necessary**

# A Program

```
int nprimes = 100;
```

data

```
int *prime, *prime2;
```

bss

```
int main() {
```

```
    int i, j, current = 1;
```

```
    prime = (int *)malloc(nprimes*sizeof(*prime));
```

```
    prime2 = (int *)malloc(nprimes*sizeof(*prime2));
```

dynamic

```
    prime[0] = 2; prime2[0] = 2*2;
```

```
    for (i=1; i<nprimes; i++) {
```

```
        NewCandidate:
```

```
        current += 2;
```

```
        for (j=0; prime2[j] <= current; j++) {
```

```
            if (current % prime[j] == 0)
```

```
                goto NewCandidate;
```

```
        }
```

```
        prime[i] = current; prime2[i] = current*current;
```

```
    }
```

```
    return 0;
```

```
}
```

text

## ... with Output

```
int nprimes = 100;
int *prime, *prime2;
int main() {
    ...
    printcol(5);
    return 0;
}
```

```
void printcol(int ncols) {
    int i, j;
    int nrows = (nprimes+ncols-1)/ncols;
    for (i = 0; i<nrows; i++) {
        for (j=0; (j<ncols) && (i+nrows*j < nvals); j++) {
            printf("%6d", prime[i + nrows*j]);
        }
        printf("\n");
    }
}
```



# ... Compiled Separately

should refer to same thing

```
int nprimes = 100;
int *prime, *prime2;
int main() {
    ...
    printcol(5);
    return 0;
}
```

**primes.c**

ditto

```
extern int nprimes;
int *prime;
void printcol(int ncols) {
    int i, j;
    int nrows = (nprimes+ncols-1)/ncols;
    for (i = 0; i<nrows; i++) {
        for (j=0; (j<ncols)
            && (i+nrows*j < nvals); j++) {
            printf("%6d", prime[i + nrows*j]);
        }
        printf("\n");
    }
}
```

**printcol.c**

**gcc -c primes.c**

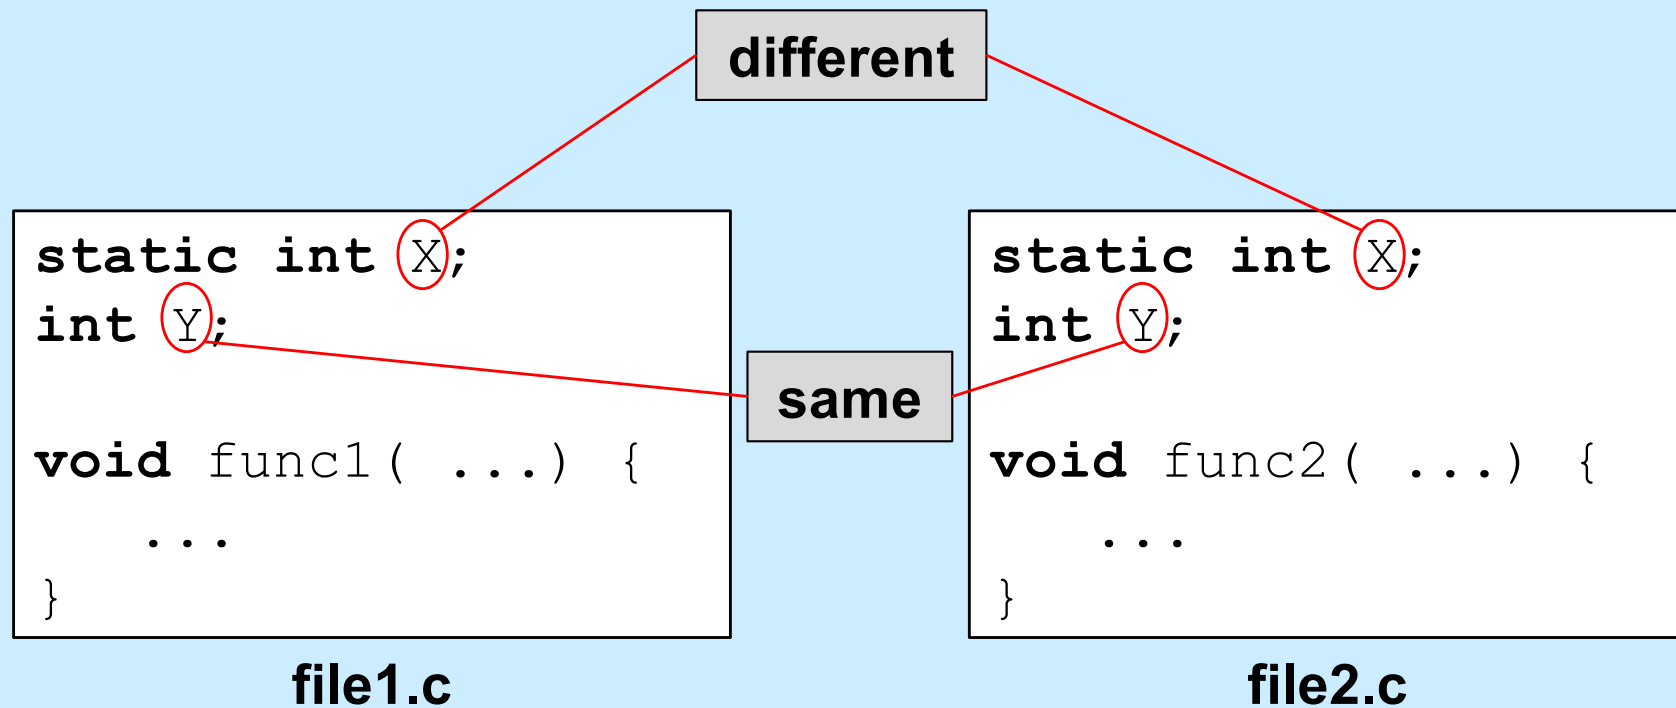
**gcc -c printcol.c**

**gcc -o primes primes.o printcol.o**

# Global Variables

- **Initialized vs. uninitialized**
  - initialized allocated in *data* section
  - uninitialized allocated in *bss* section
    - » implicitly initialized to zero
- **File scope vs. program scope**
  - *static* global variables known only within file that declares them
    - » two of same name in different files are different
    - » e.g., `static int X;`
  - non-static global variables potentially shared across all files
    - » two of same name in different files are same
    - » e.g., `int X;`

# Scope



# Static Local Variables

```
int *sub1() {  
    int var = 1;  
    ...  
    return &var;  
    /* amazingly illegal */  
}
```

```
int *sub2() {  
    static int var = 1;  
    ...  
    return &var;  
    /* (amazingly) legal */  
}
```

# Reconciling Program Scope (1)

tentative definition

```
int X;  
  
void func1( ...) {  
    ...  
}
```

file1.c

(complete) definition

```
int X=1;  
  
void func2( ...) {  
    ...  
}
```

file2.c

**Where does X go?**  
**What's its initial value?**

- tentative definitions overridden by compatible (complete) definitions
- if not overridden, then initial value is zero

# Reconciling Program Scope (2)

```
int X=2;  
  
void func1( ...) {  
    ...  
}
```

**file1.c**

```
int X=1;  
  
void func2( ...) {  
    ...  
}
```

**file2.c**

**What happens here?**

# Reconciling Program Scope (3)

```
int X=1;

void func1( ...) {
    ...
}
```

**file1.c**

```
int X=1;

void func2( ...) {
    ...
}
```

**file2.c**

**Is this ok?**

# Reconciling Program Scope (4)

```
extern int X;  
  
void func1( ...) {  
    ...  
}
```

file1.c

```
int X=1;  
  
void func2( ...) {  
    ...  
}
```

file2.c

What's the purpose of “extern”?



# Does Location Matter?

```
int main(int argc, char *[]) {  
    return(argc);  
}
```

main:

```
pushq %rbp    ; push frame pointer  
movq  %rsp, %rbp    ; set frame pointer to point to new frame  
movl  %edi, %eax    ; put argc into return register (eax)  
movq  %rbp, %rsp    ; restore stack pointer  
popq  %rbp    ; pop stack into frame pointer  
ret         ; return: pops end of stack into rip
```

# Location Matters ...

```
int X=6;  
int *aX = &X;
```

```
int main() {  
    void subr(int);  
    int y=*aX;  
    subr(y);  
    return(0);  
}
```

```
void subr(int i) {  
    printf("i = %d\n", i);  
}
```

# Coping

- **Relocation**

- modify internal references according to where module is loaded in memory
- modules needing relocation are said to be *relocatable*
  - » which means they *require* relocation
- the compiler/assembler provides instructions to the linker on how to do this